

Image Warping Report - Lab 2

What I implemented & How

we focus on the classes and methods I wrote instead of the algorithms that are explained in the homework documents.

I implemented the following three buttons under the "Warping" menu bar.

IDW (Inverse distance weighted)

- I implemented *local_weight_function* which is then used for *weight_function* here I chose the other weight function presented in the paper with $R = 2500$ and $\mu = 3$, then I implemented *local_appro_with_matrix* which is used for *local_appro*, it takes in a vector of matrices D_i , here D_i 's are simply assumed to be zero.
- Then I combined the two and got a function *transform* which is then used in the main *warp* function. If D_i 's are optimized then we can feed them back to *local_appro_with_matrix* and everything works out. I implemented an optimization process in *initialize()* but there are bugs I haven't fixed yet - it runs but somehow the gradient is too small, thus the optimization process is not started. I intended to use the *LevenbergMarquardt* algorithm.

RBF (Radial basis function)

First I implemented *Gaussian()* with $\sigma = 200$, here I picked the transition function in the paper instead of Gaussian. I maintain an `Eigen::MatrixXf` private to keep track of the weights.

Then in *initialize()*, we obtain the weights in this function by solving the linear equation, which is then used in a private helper function *transform()*.

Gap filling (used with RBF)



I defined a structure called *Pixel* with three fields. Then maintain two vectors of Pixels: *mapped_to* and *unmapped_to*.

gap_filling()

- Our helper function *gap_filling* takes in these two vectors and returns another vector of Pixels - those Pixels that were not mapped to with their interpolated values.
- I build an `AnnoyIndex<int, float, Euclidean, Kiss32Random, AnnoyIndexSingleThreadedBuildPolicy>` in *gap_filling*, find its k nearest neighbors, and take the distance-weighted average of their values, $k = 70$ in our case.
- To discern the boundary, I count the number of uncolored pixels in its k neighbors, only when $k \leq 1$ do we update its pixel value, this way we can detect the boundary when contracting the picture. One can tune these parameters to adjust how blurred-out the picture is: k , uncolored count, weights used for average_value.
- It takes roughly 30 seconds to execute one warping.



Result



IDW







IDW behaves poorly when the number of points is small because D_i 's are not optimized.

RBF














RBF with gap-filling







Future Work (Features I haven't implemented)

For gap-filling

- If we implement a boundary detect method in the *warp()* for gap filling, I've tried that it can deal with the contraction cases perfectly, but not with expansion since it neglects too many uncolored pixels. Thus I deleted it, but maybe we can make it work by tuning its parameters.
- tune parameters in *gap_filling()*.
- optimize the algorithm for better performance and shorter running time.

For IDW optimization

- there is something wrong with the setup of `struct IDW::ObjectiveFunction : Functor < float >`, the method runs but the optimization process isn't started due to zero gradient. We likely need to adjust how we set up the optimizer, such as its *df* field.
- maybe change the optimization algorithm, using TensorFlow instead.
- Tune the parameters of the weight function in IDW.cpp